



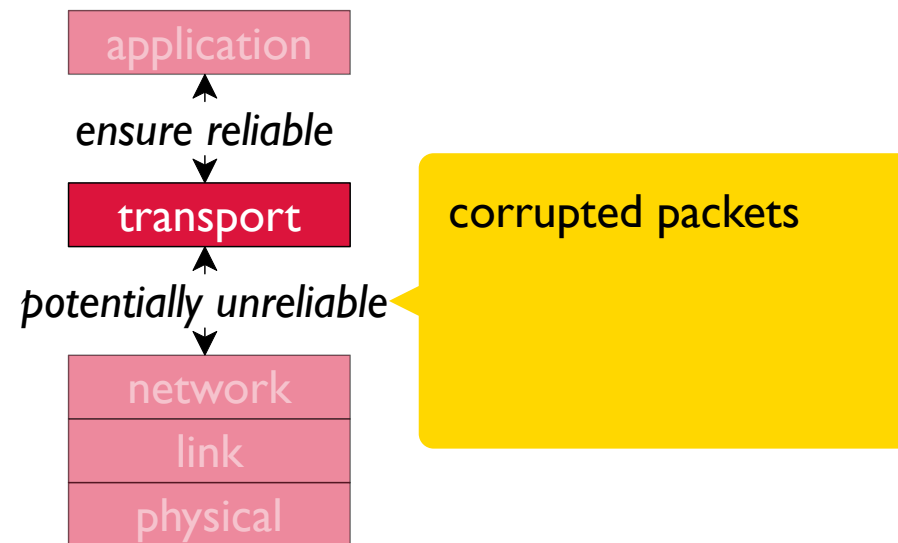
Reliable Data Transfer



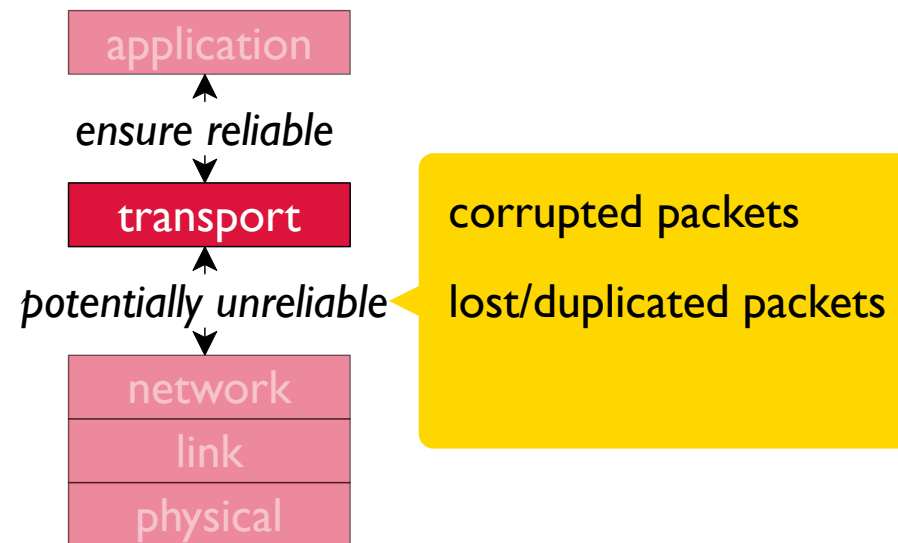
Reliable Data Transfer



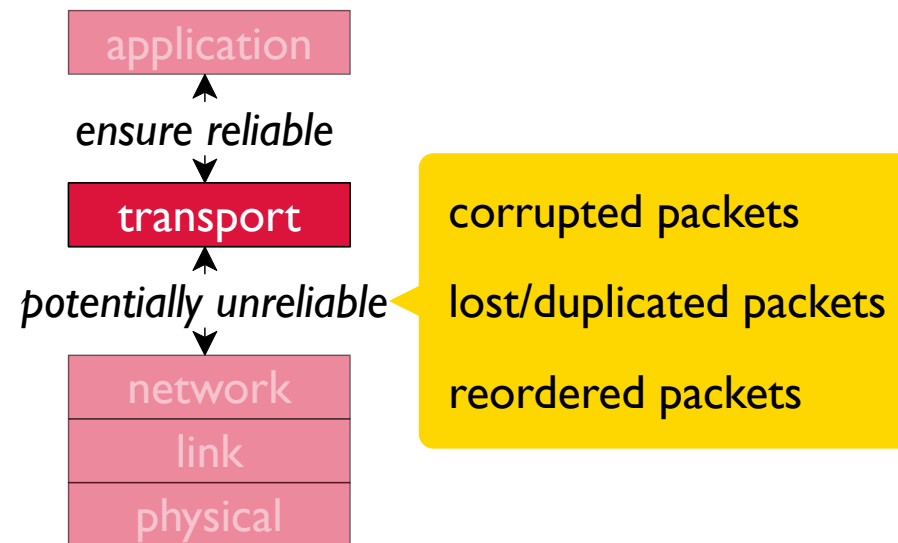
Reliable Data Transfer




Reliable Data Transfer



Reliable Data Transfer




Reliable Data Transfer




Reliable Data Transfer

“reliable data transfer”




Reliable Data Transfer

“reliable data transfer”




“unreliable data transfer”


Reliable Data Transfer




Reliable Data Transfer



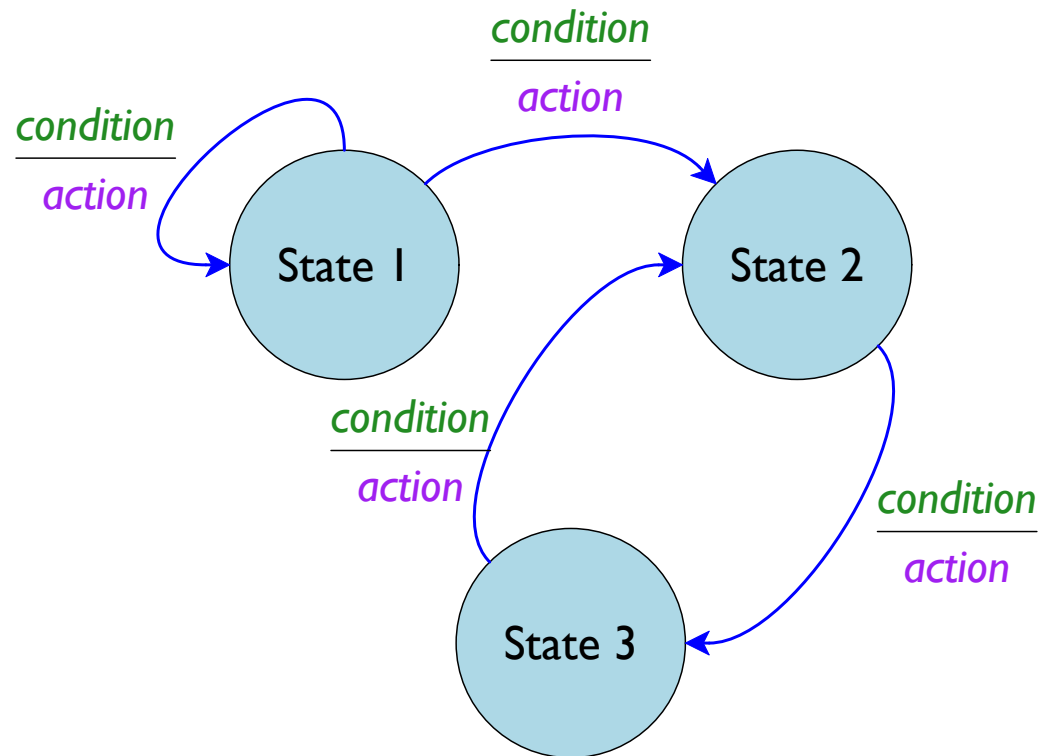
Reliable Data Transfer



State Machines

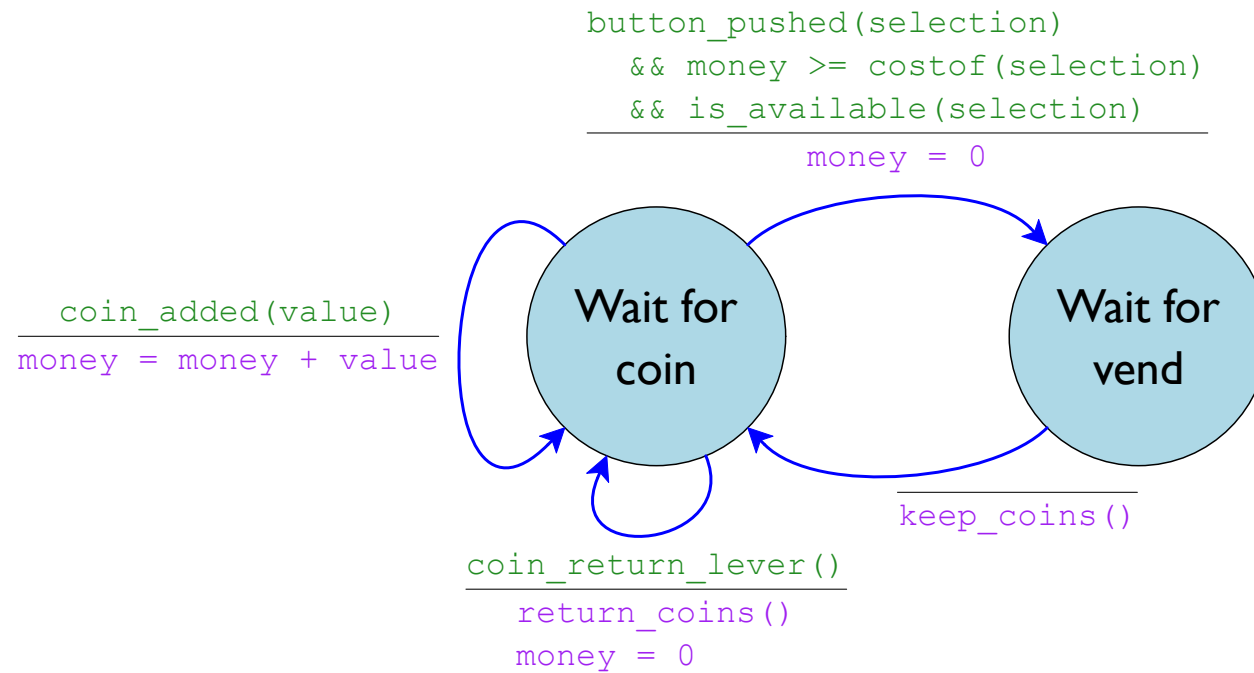


State Machines



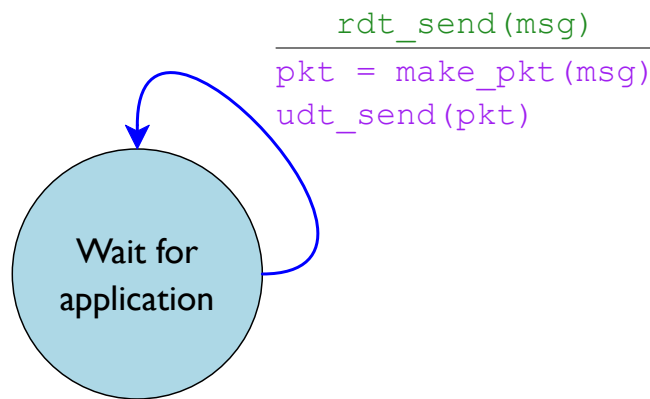
State Machines

Vending machine example:



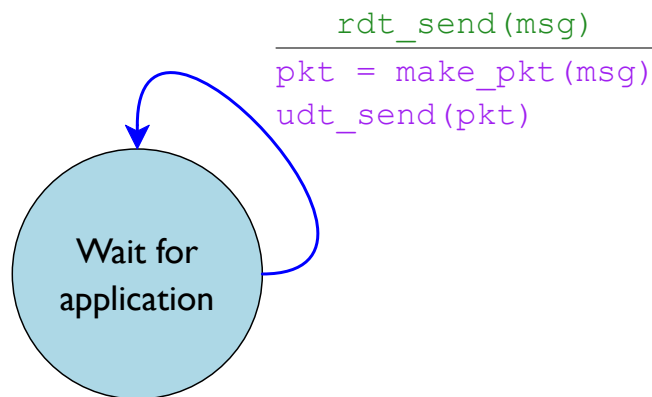
Assuming Reliable `udt_send` and `udt_rcv`

sending host

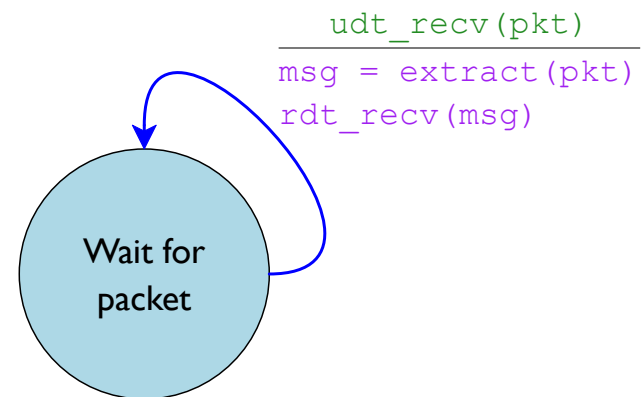


Assuming Reliable `udt_send` and `udt_rcv`

sending host



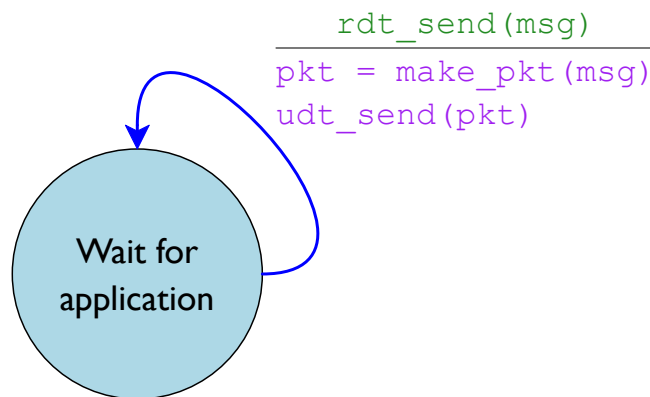
receiving host



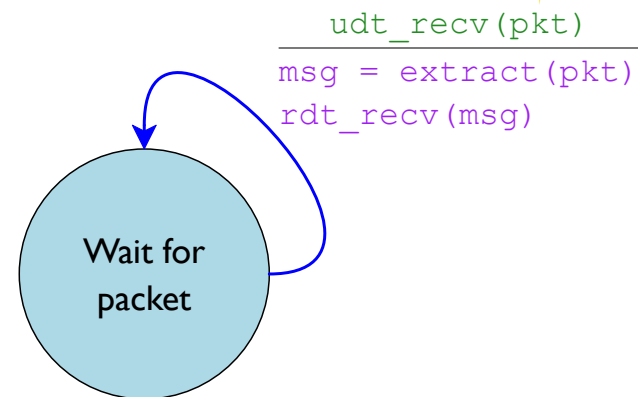
Assuming Reliable `udt_send` and `udt_rcv`

What if `pkt` is corrupted?


sending host



receiving host




Checksum




Checksum

I'd like 1 apple, 2 bananas, and 3 cherries



Checksum


I'd like 1 apple, 2 bananas, and 3 cherries



Ok: 1 apple, 2 bananas, and 2 cherries


Checksum

I'd like 1 apple, 2 bananas, and 3 cherries — which is 6 total



Checksum


I'd like 1 apple, 2 bananas, and 3 cherries — which is 6 total



Ok: 1 apple, 2 bananas, and 2 cherries — which is 6... oops

Checksum

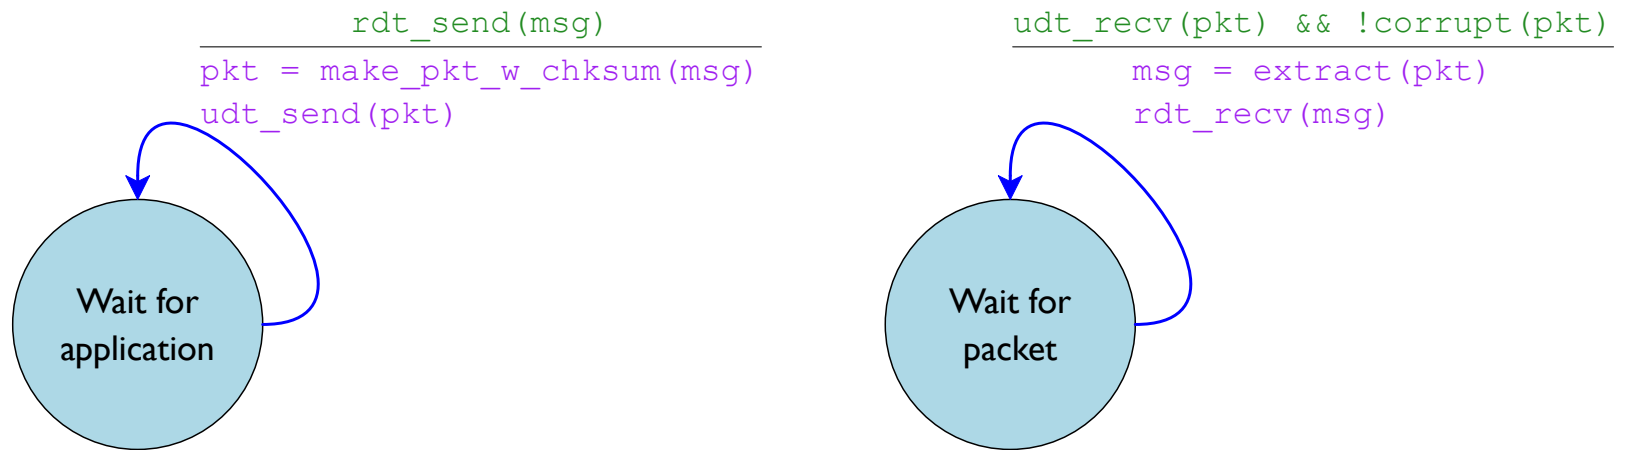
I'd like 1 apple, 2 bananas, and 3 cherries — which is 6 total



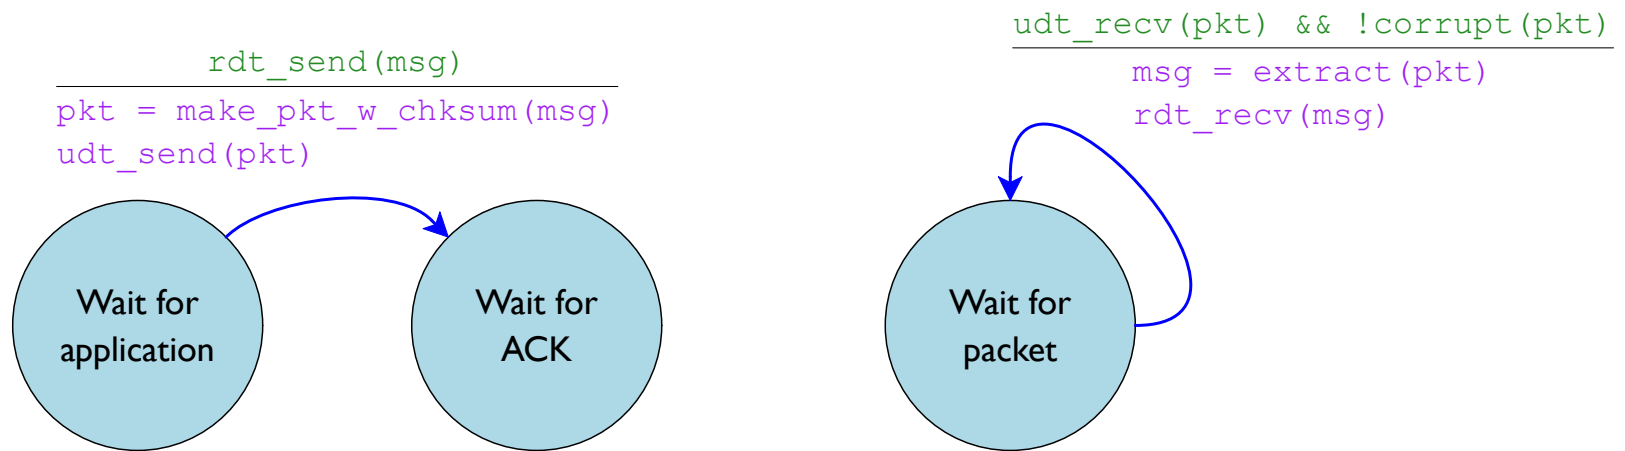
Ok: 1 apple, 2 bananas, and 2 cherries — which is 6... oops

To deal with lots of numbers, just keep low bits of the sum

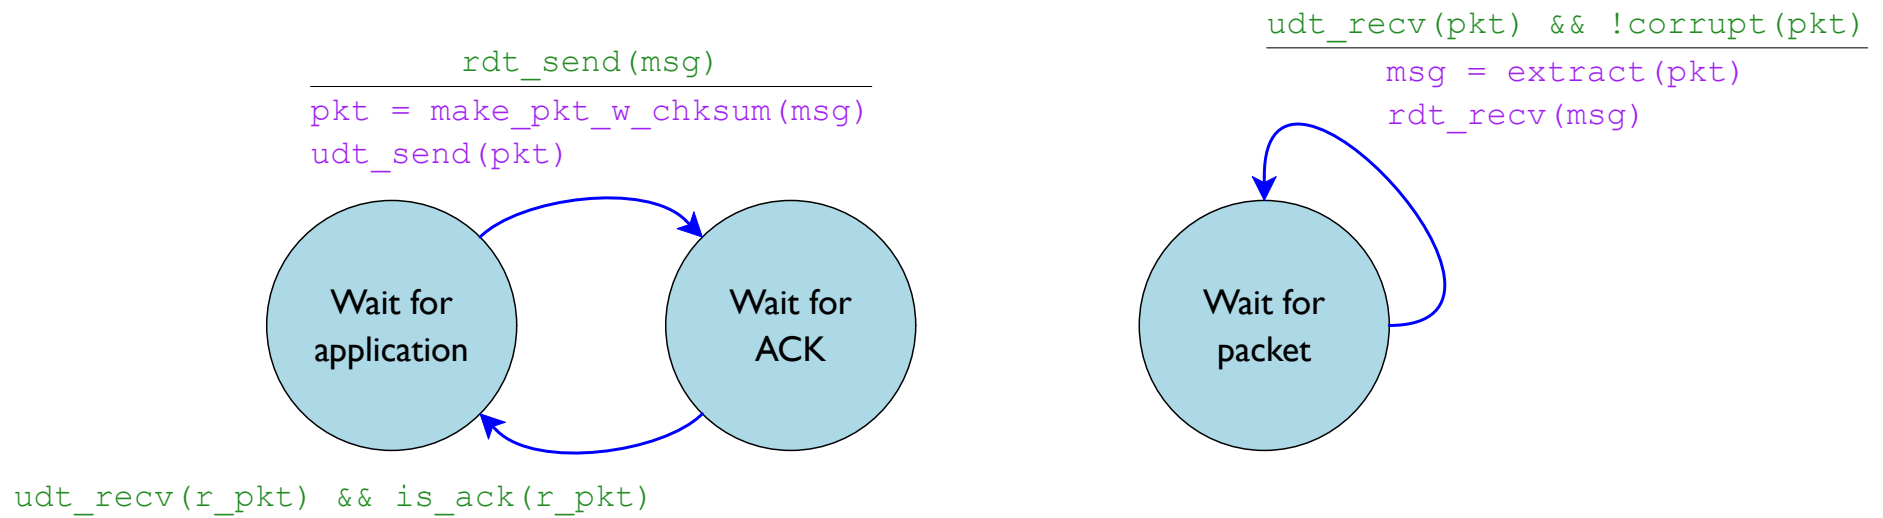
Using a Checksum



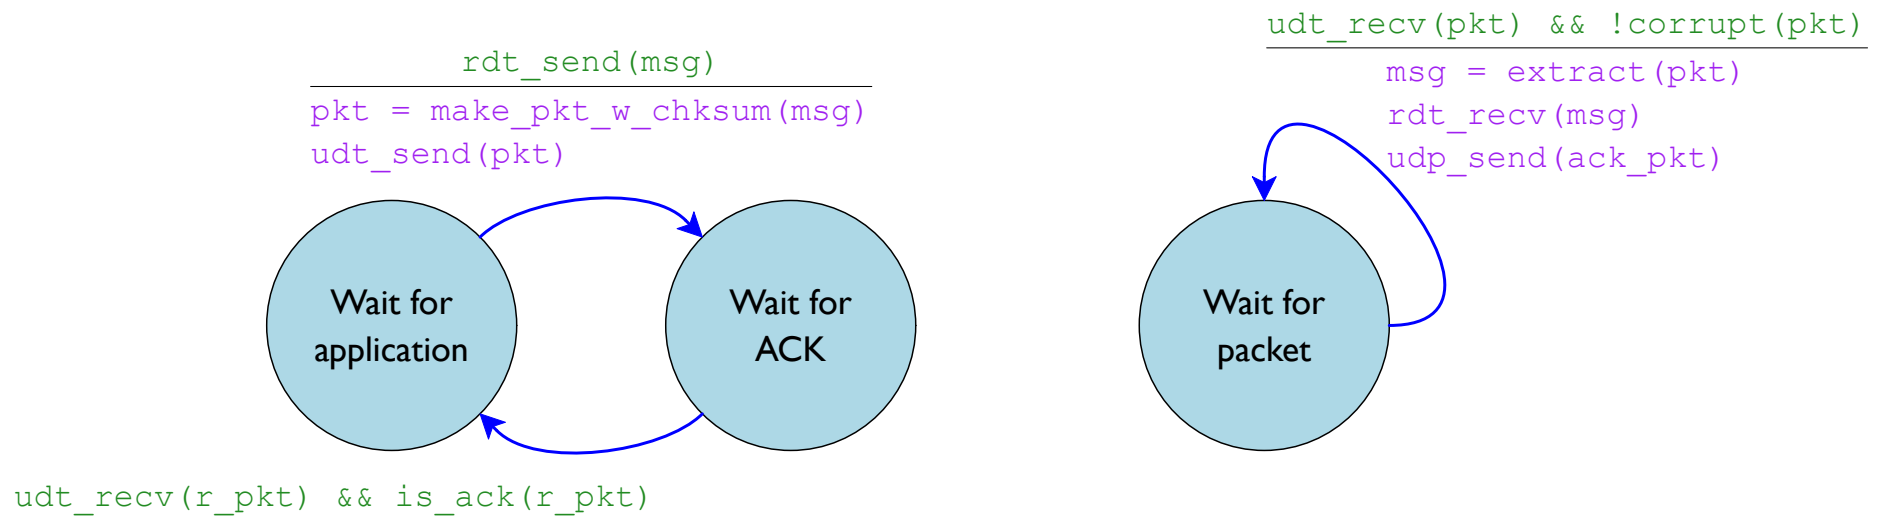
Using a Checksum



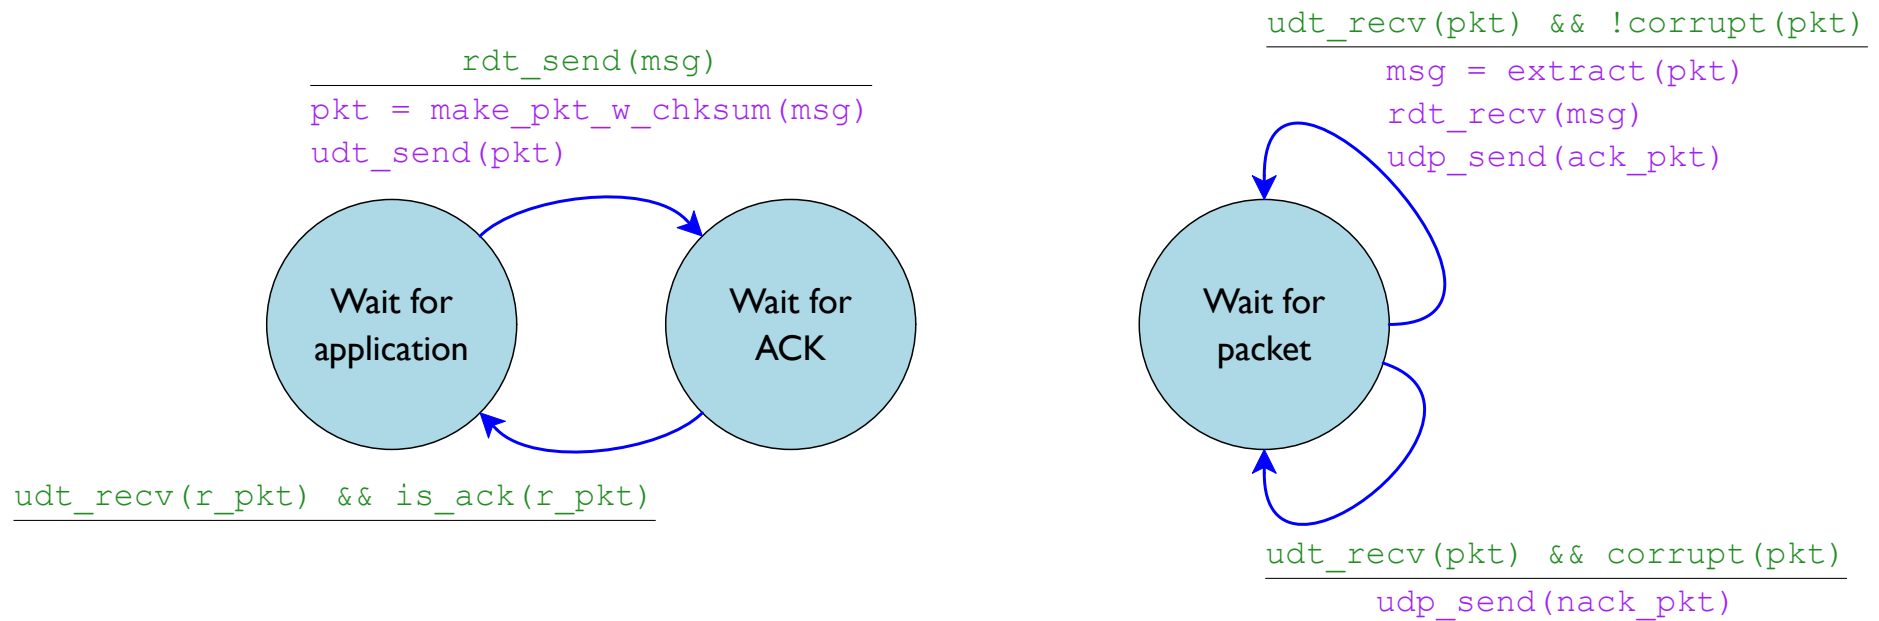
Using a Checksum



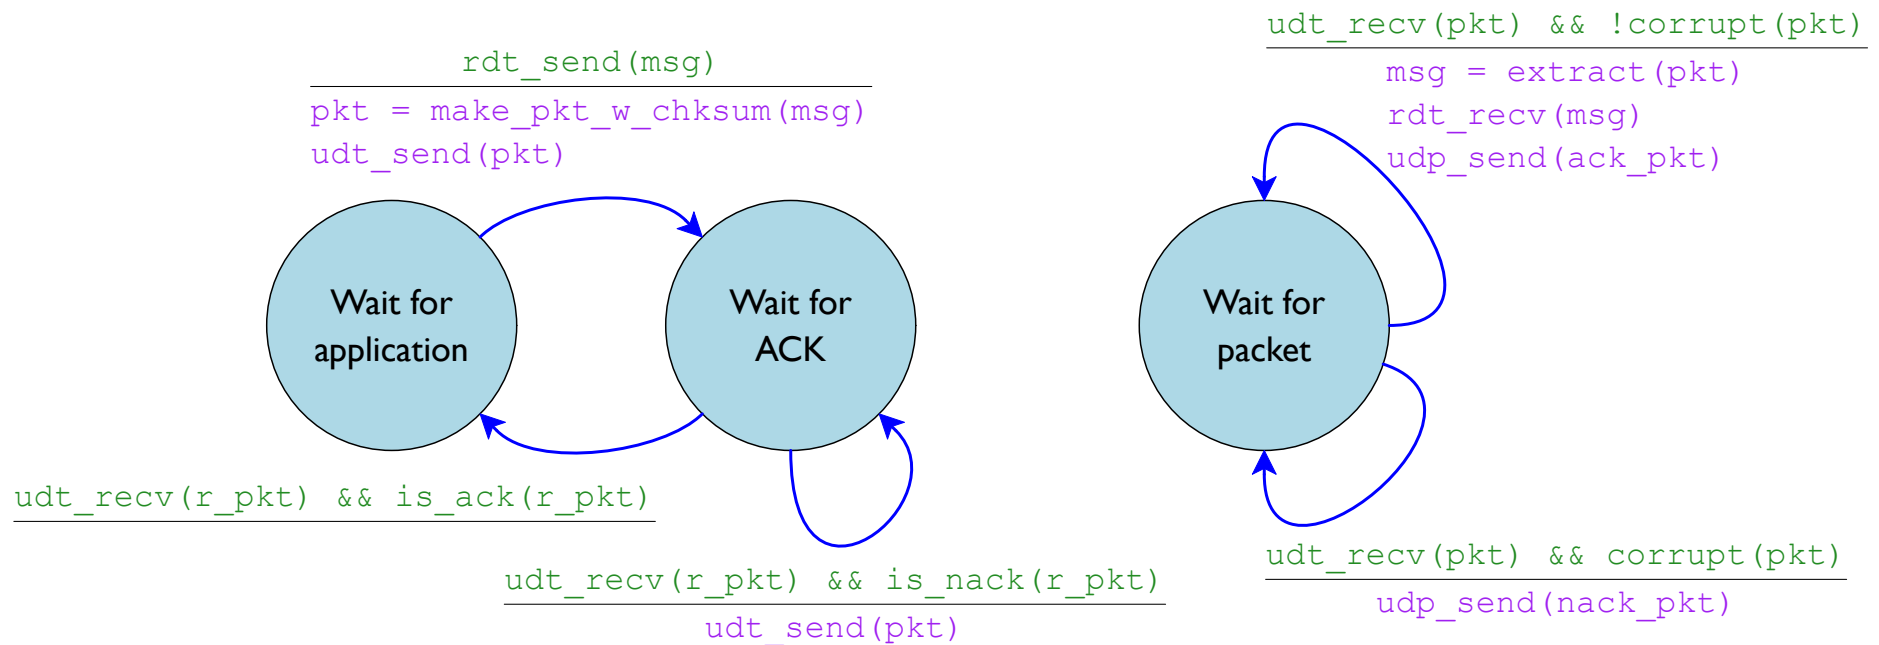
Using a Checksum



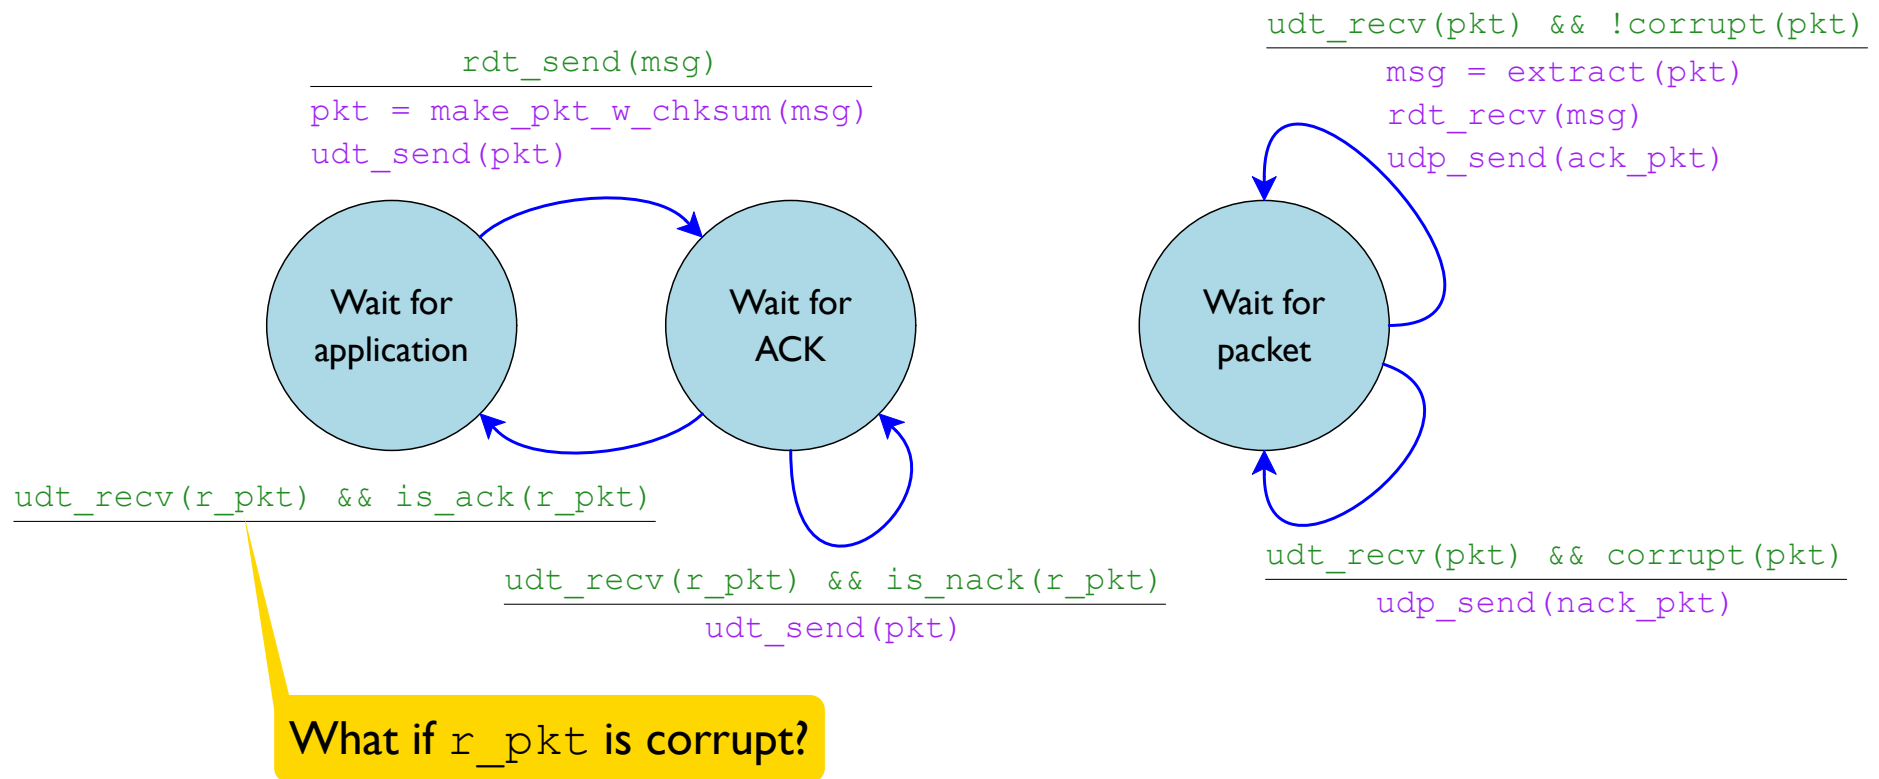
Using a Checksum



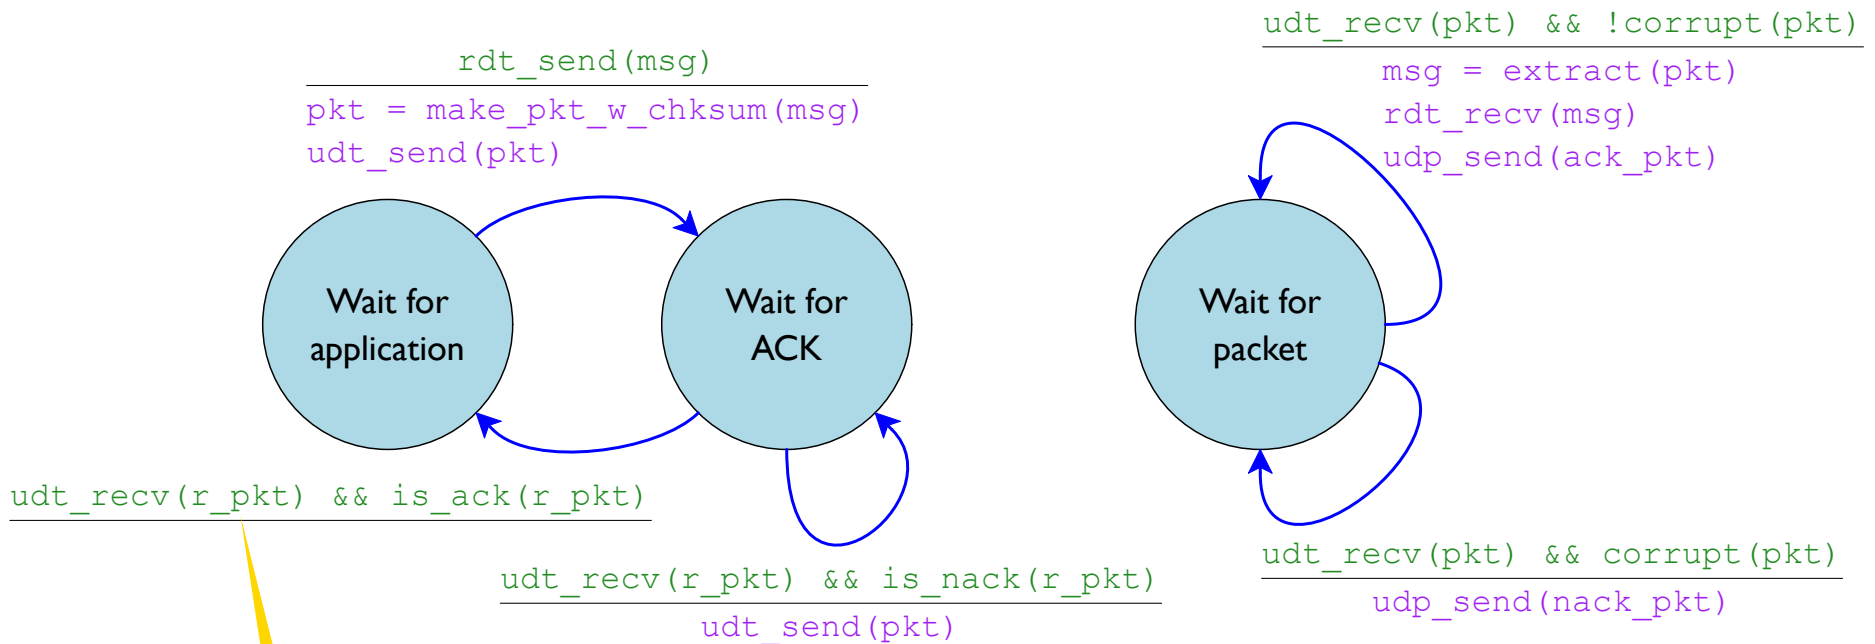
Using a Checksum



Using a Checksum

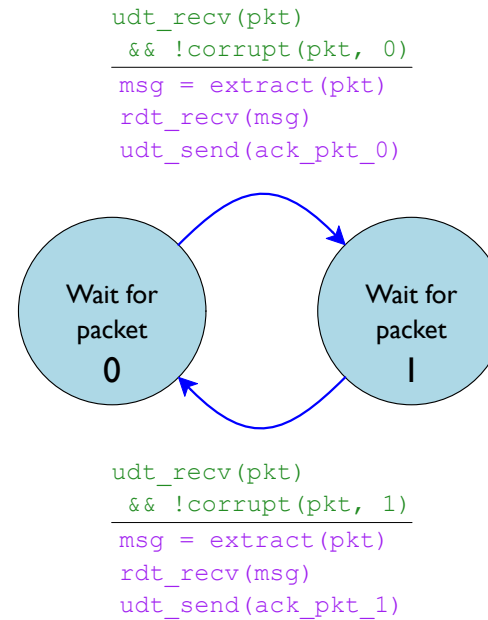
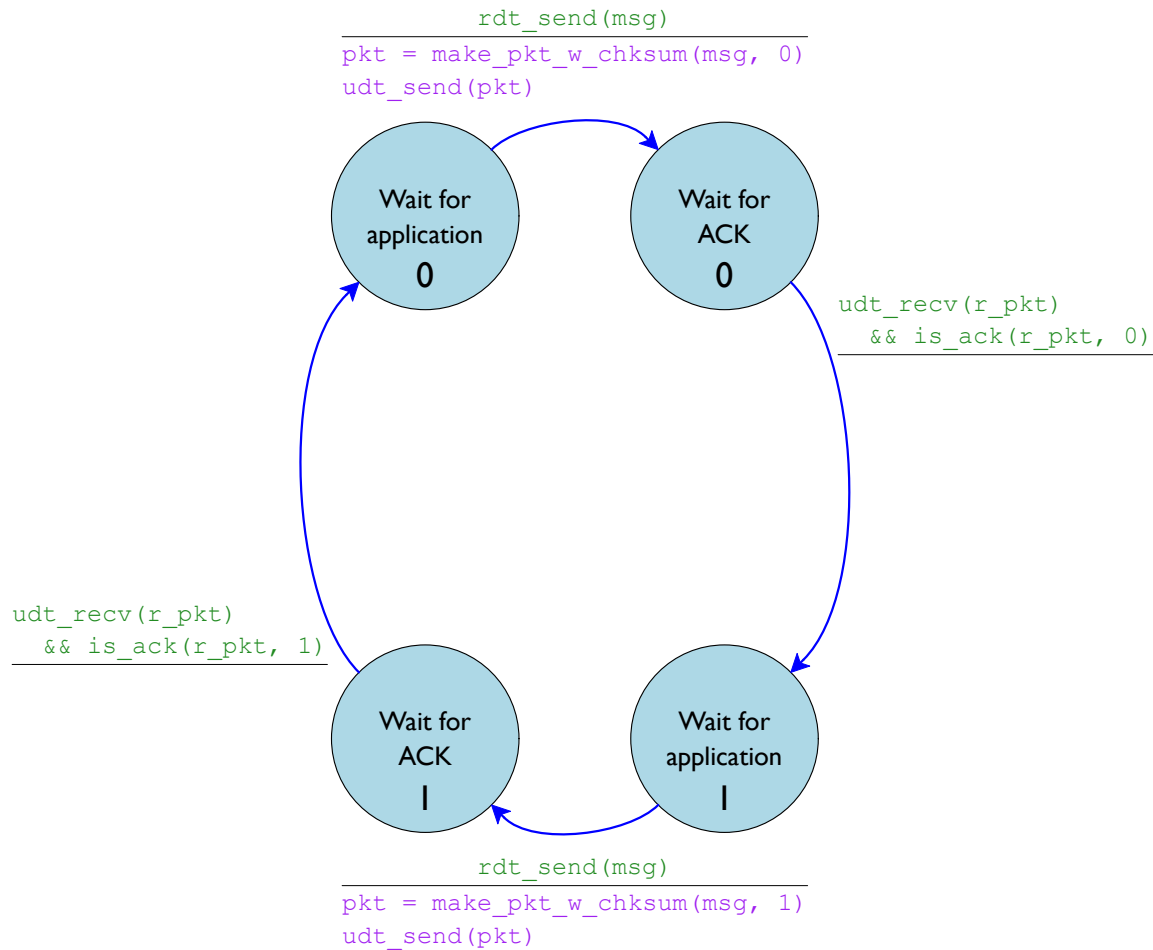


Using a Checksum

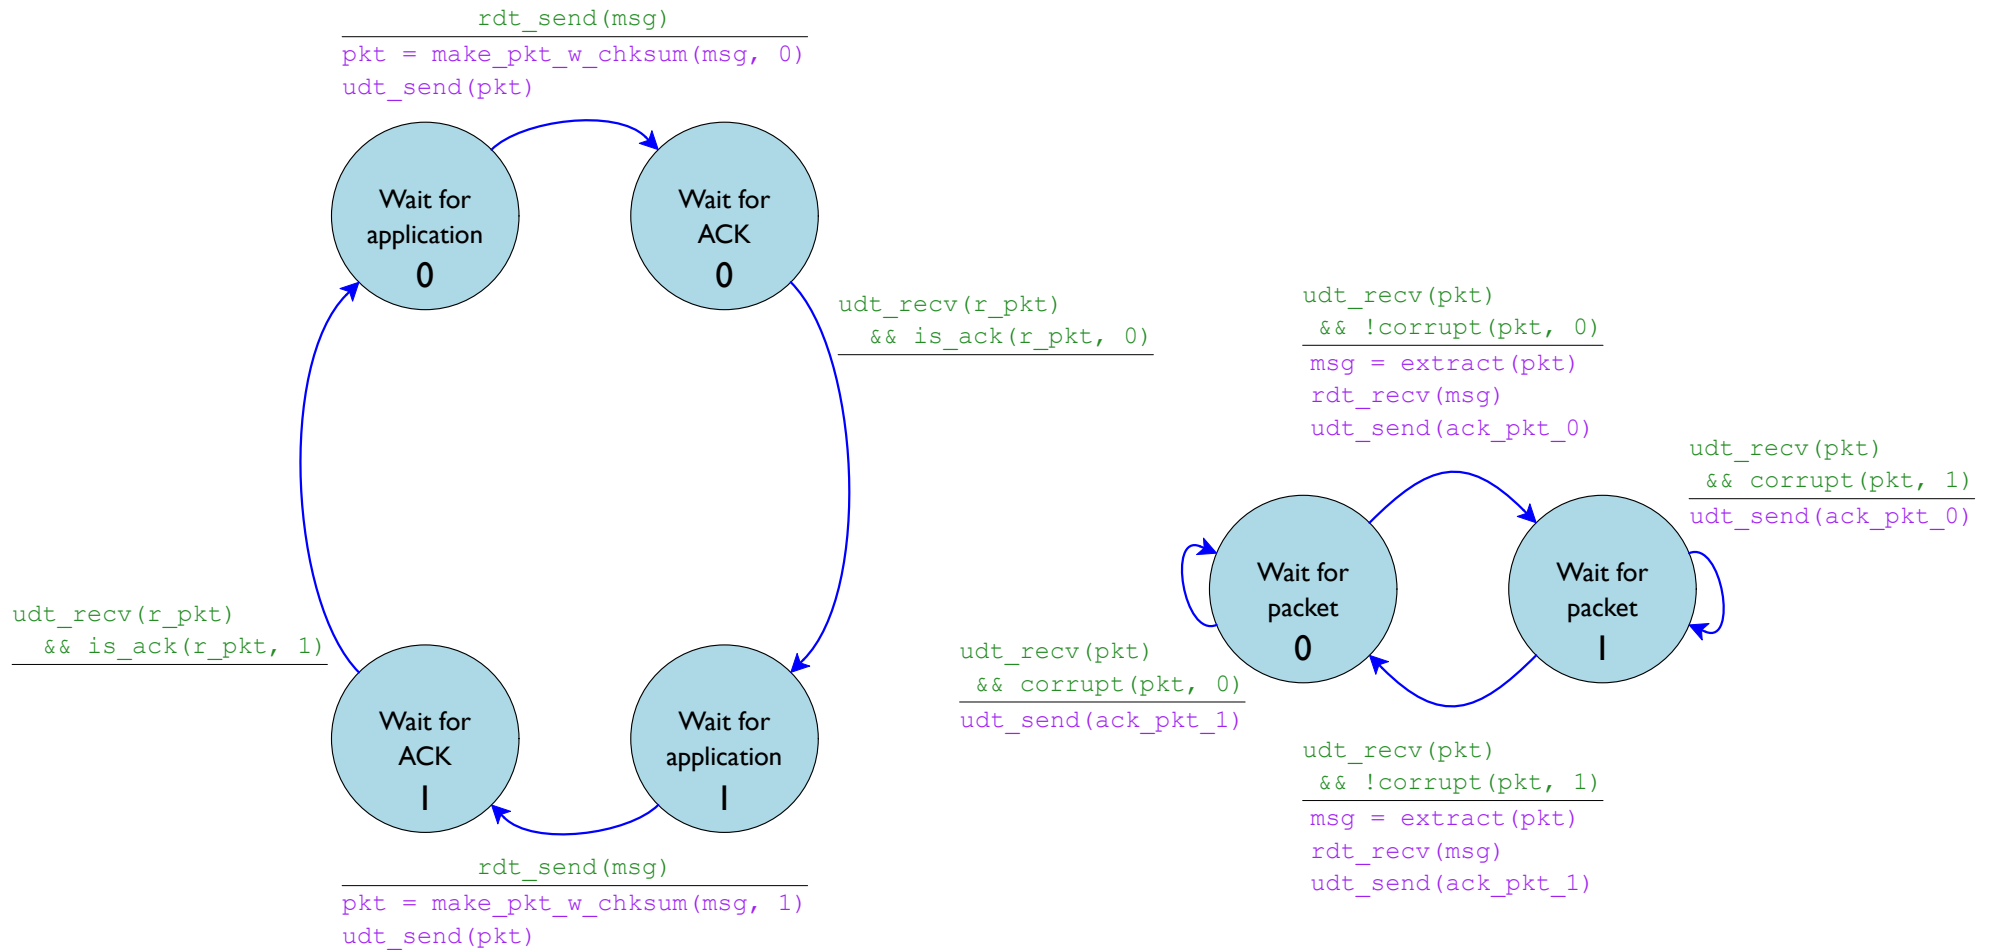


How do we avoid an ACK of ACK of ACK...?

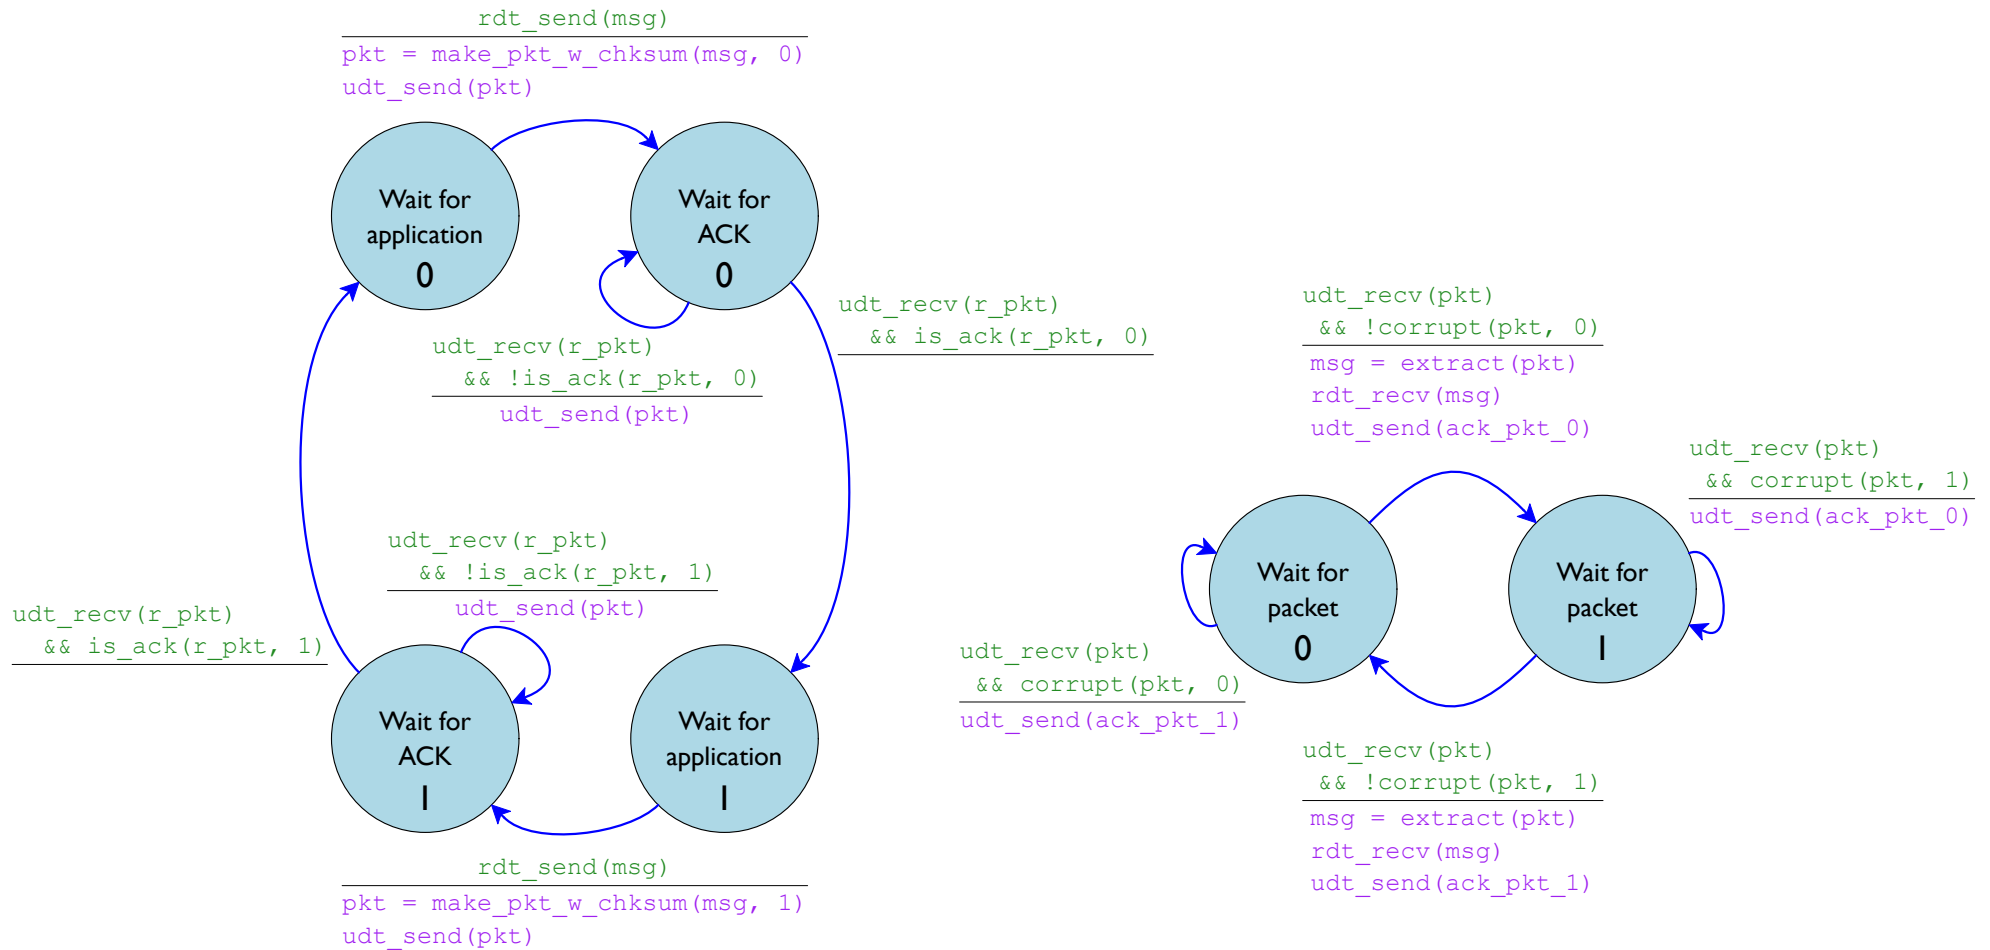
Handling ACK Corruption



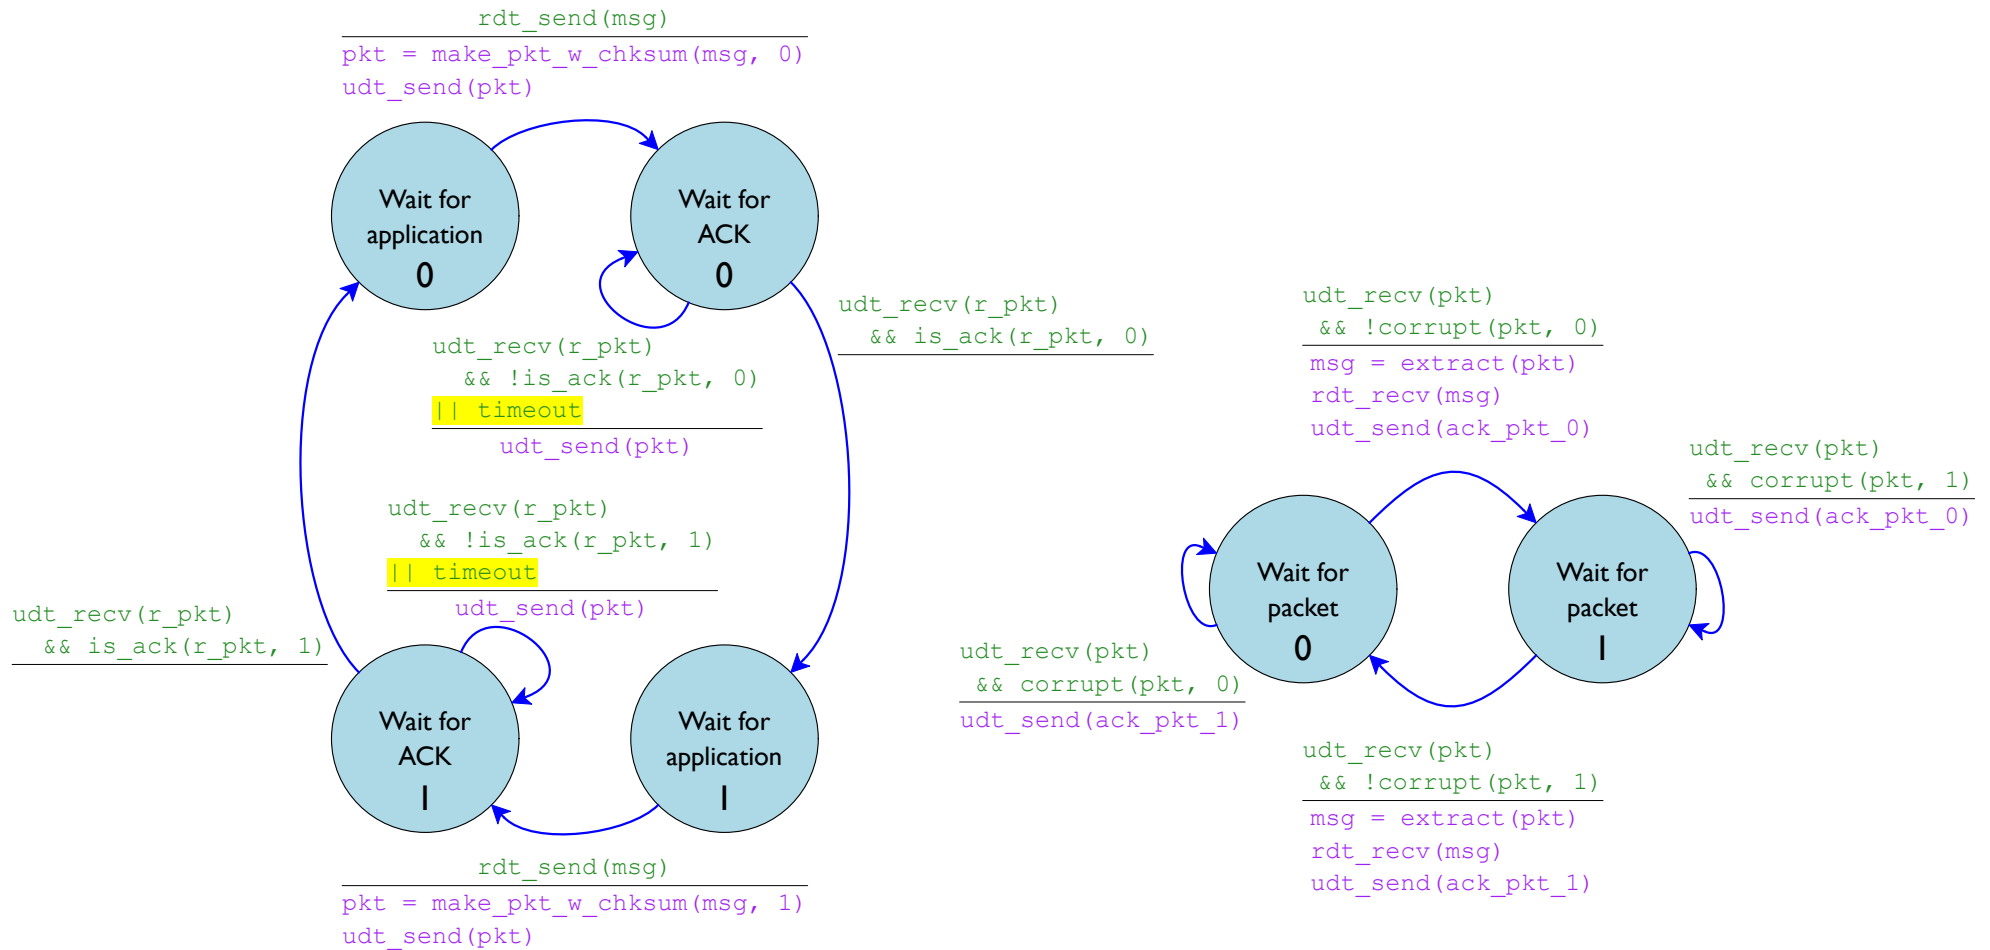
Handling ACK Corruption



Handling ACK Corruption

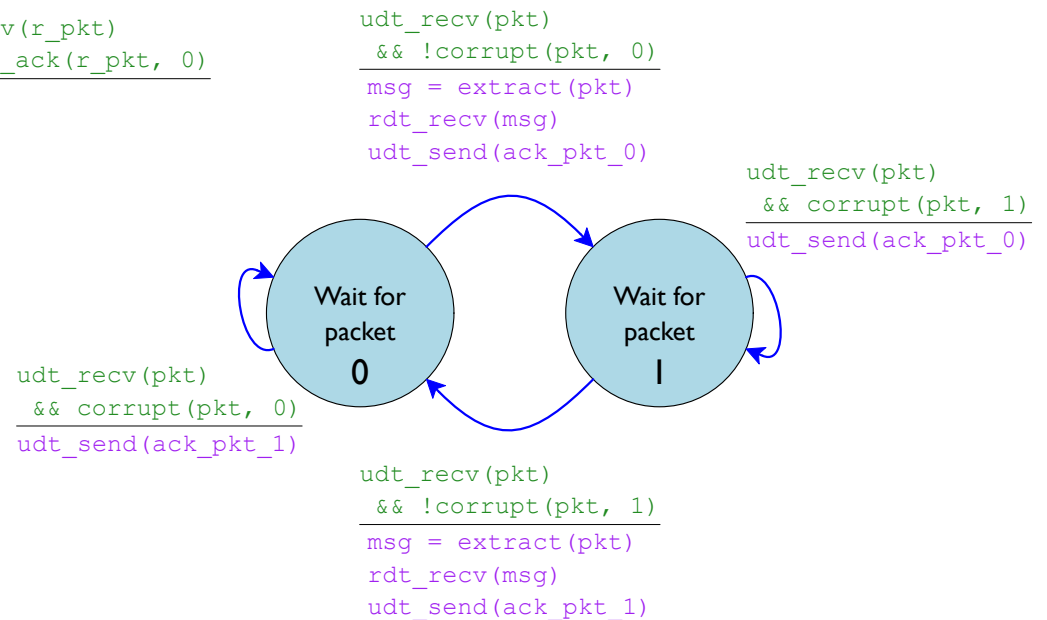
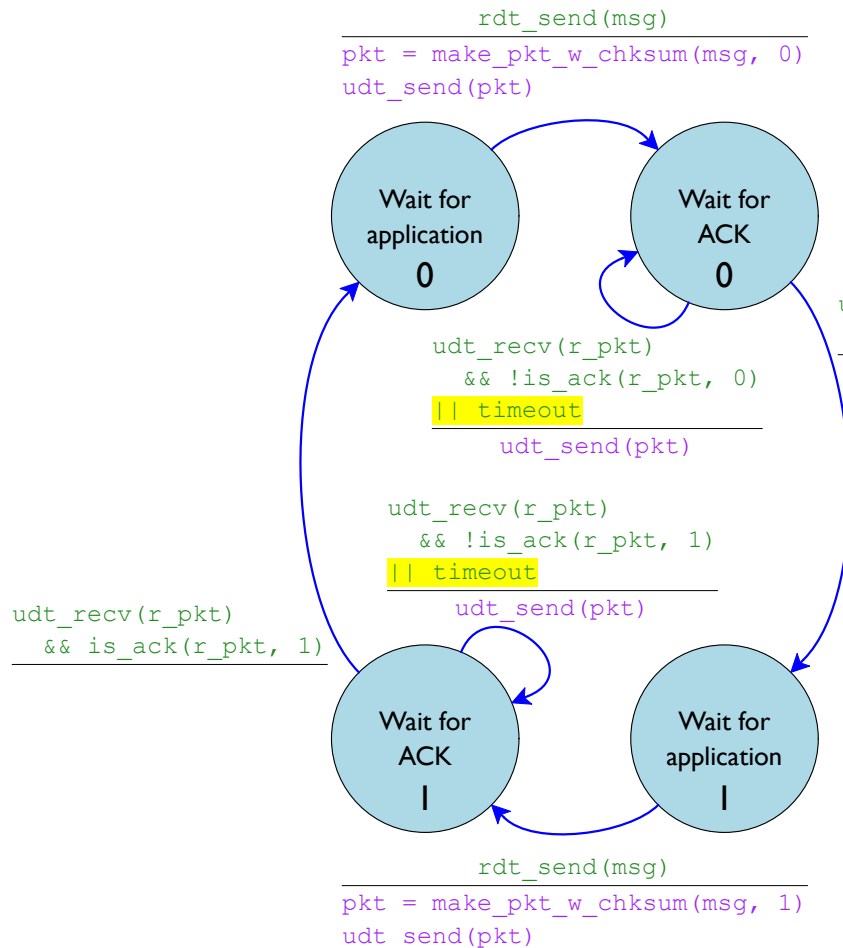


Handling Corrupt and Lost Packets



Handling Corrupt and Lost Packets

To handle duplicated and reordered packets, use a sequence number that always counts up instead of just 0 and 1



Choosing a Timeout

RTT is minimum useful timeout

Choosing a Timeout

RTT is minimum useful timeout

- too small \Rightarrow resend data and ACKs unnecessarily
- too large \Rightarrow sender waits too long to resend

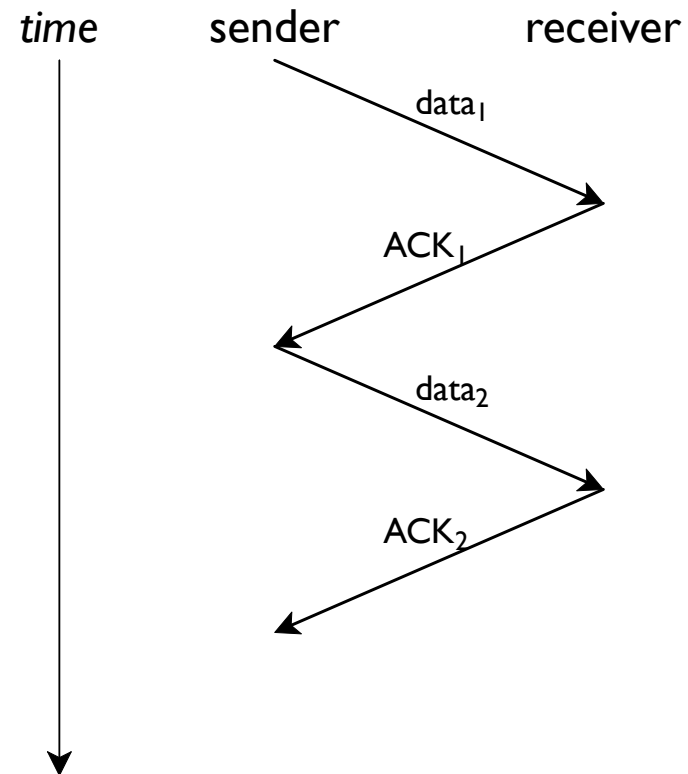
Choosing a Timeout

RTT is minimum useful timeout


- too small \Rightarrow resend data and ACKs unnecessarily
- too large \Rightarrow sender waits too long to resend

$scale \times avg(RTT) + stddev(RTT)$ is a good approach

Sequential Messages




Sequential Messages




Throughput is limited by latency

Pipelined Messages



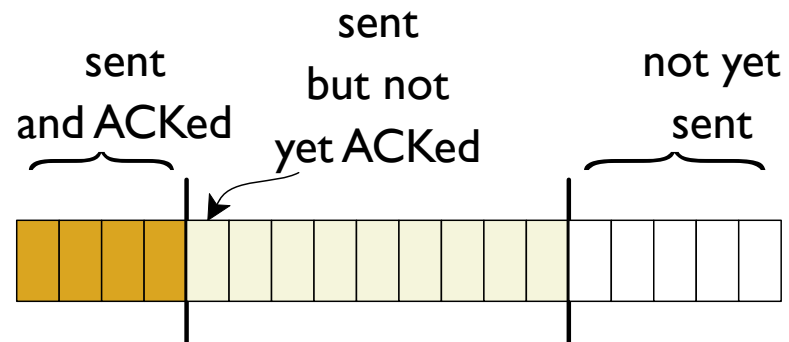
Pipelined Messages




Need a way to track multiple messages in flight

Buffers

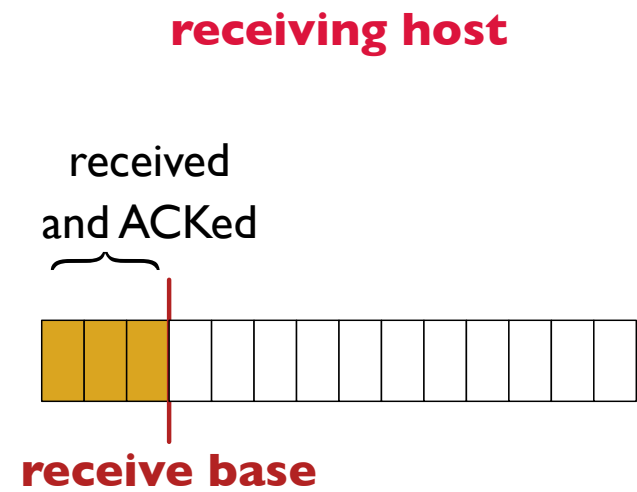
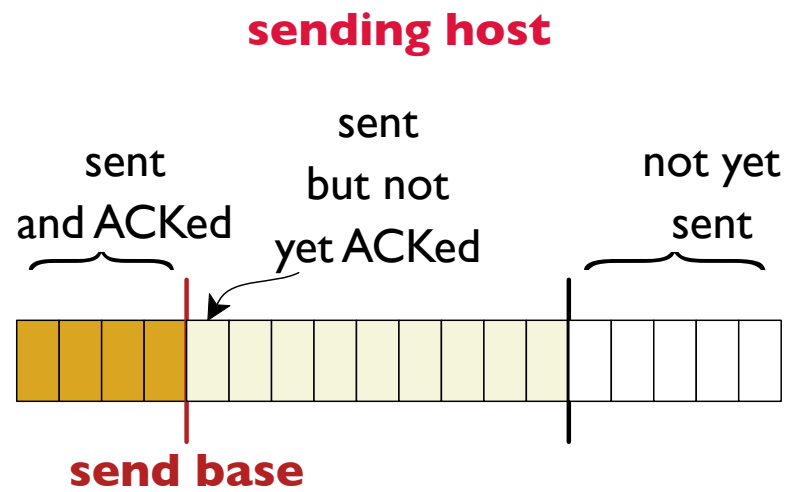
sending host



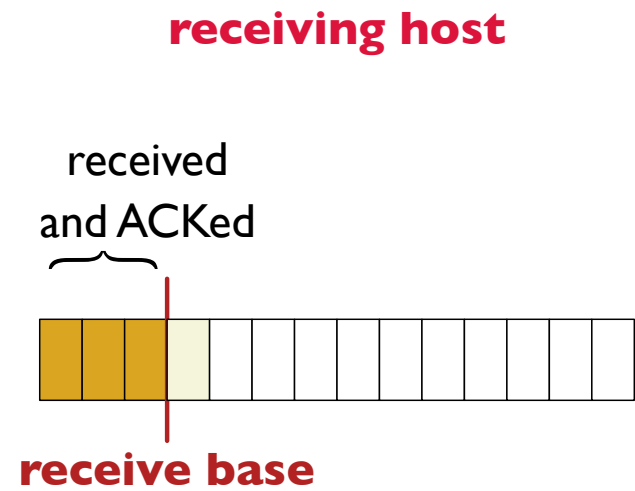
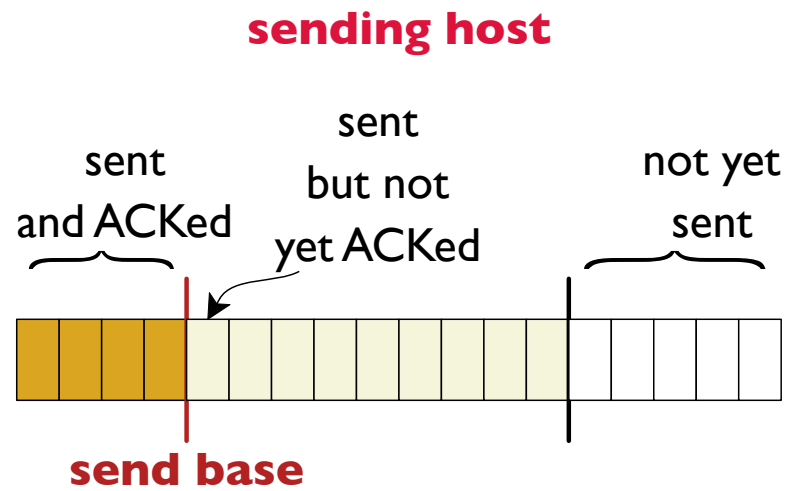
receiving host



Buffers

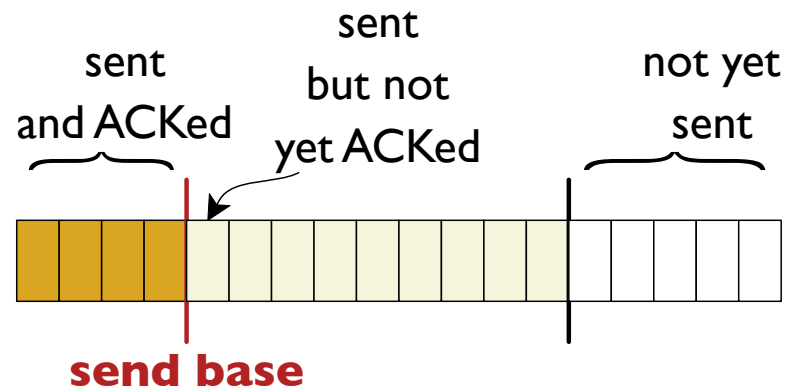


Buffers




Buffers

sending host

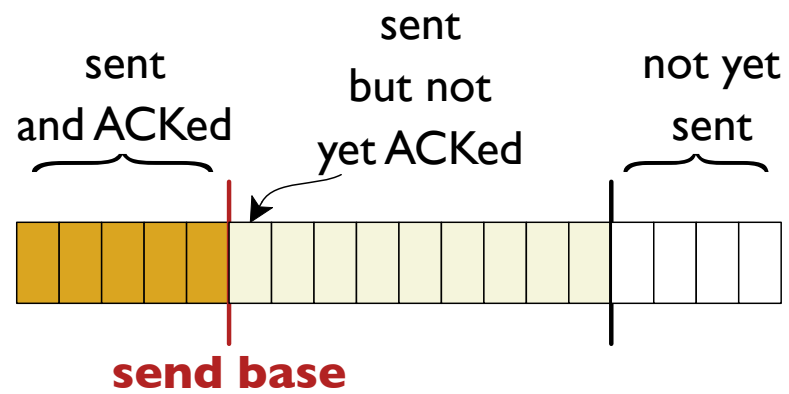


receiving host




Buffers

sending host




receiving host




Buffers

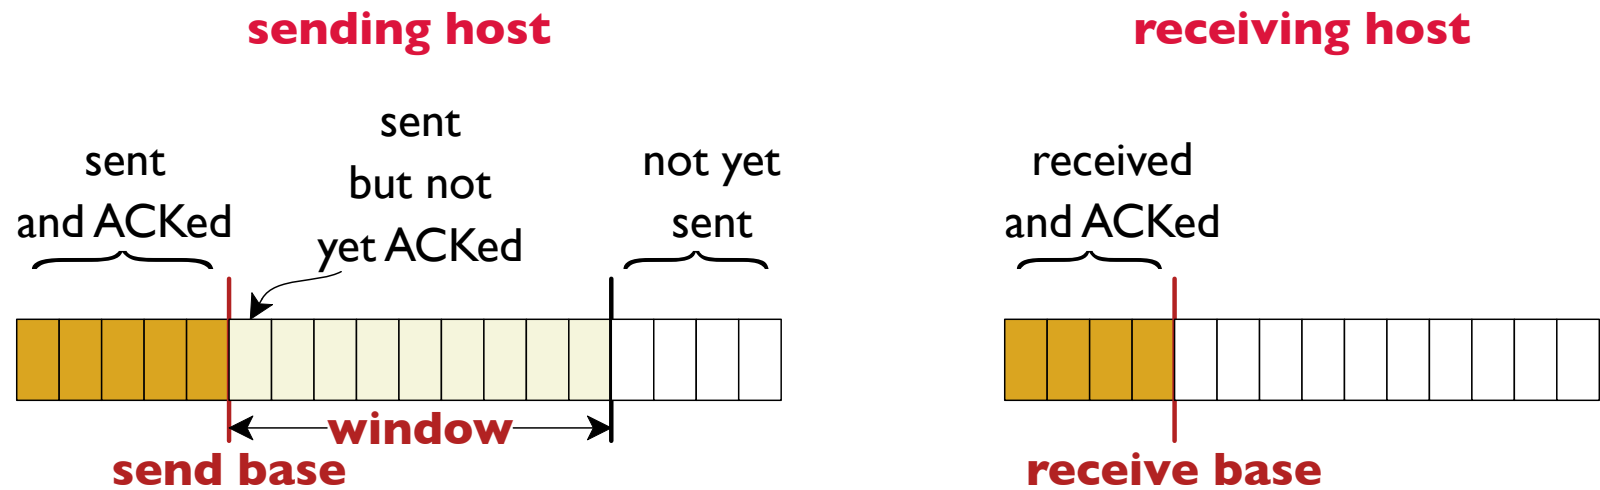
sending host



receiving host




Buffers




Like a timeout, the window size needs to be chosen well

Buffers

sending host




receiving host




Buffers

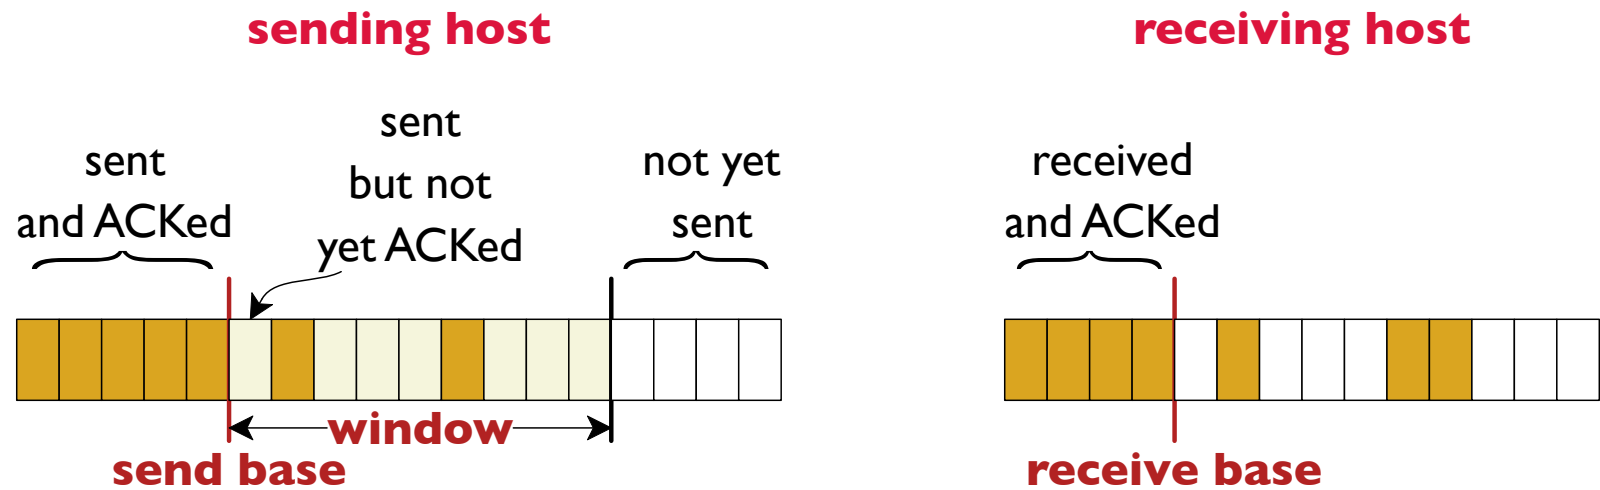
sending host



receiving host

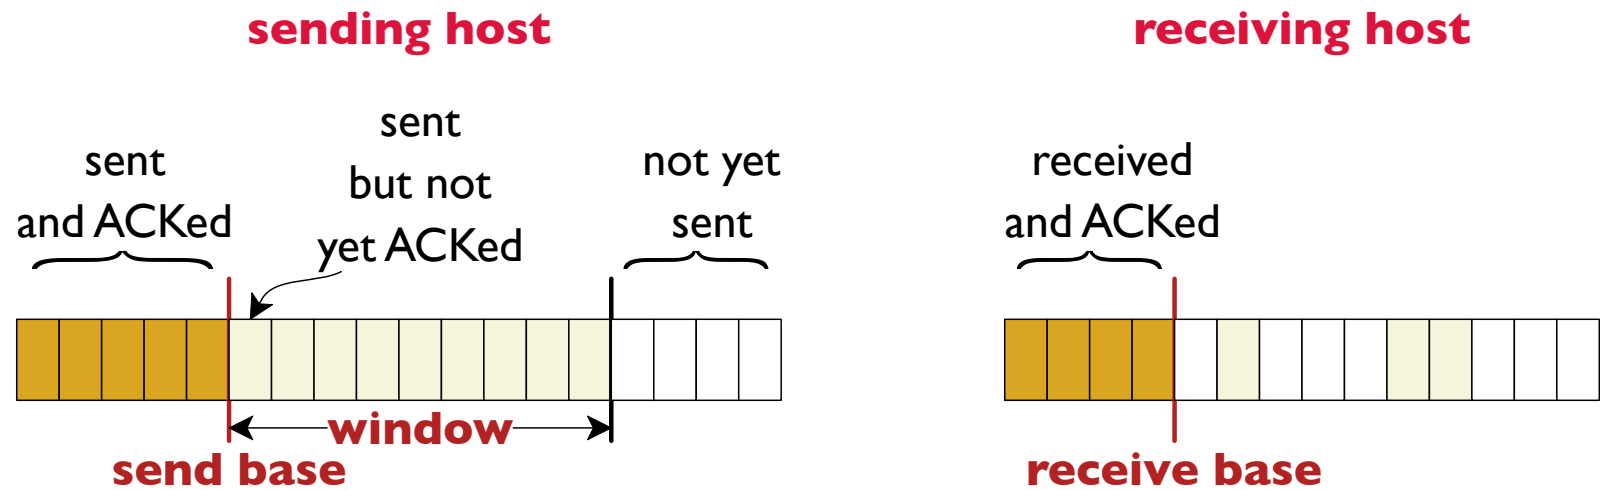


Buffers



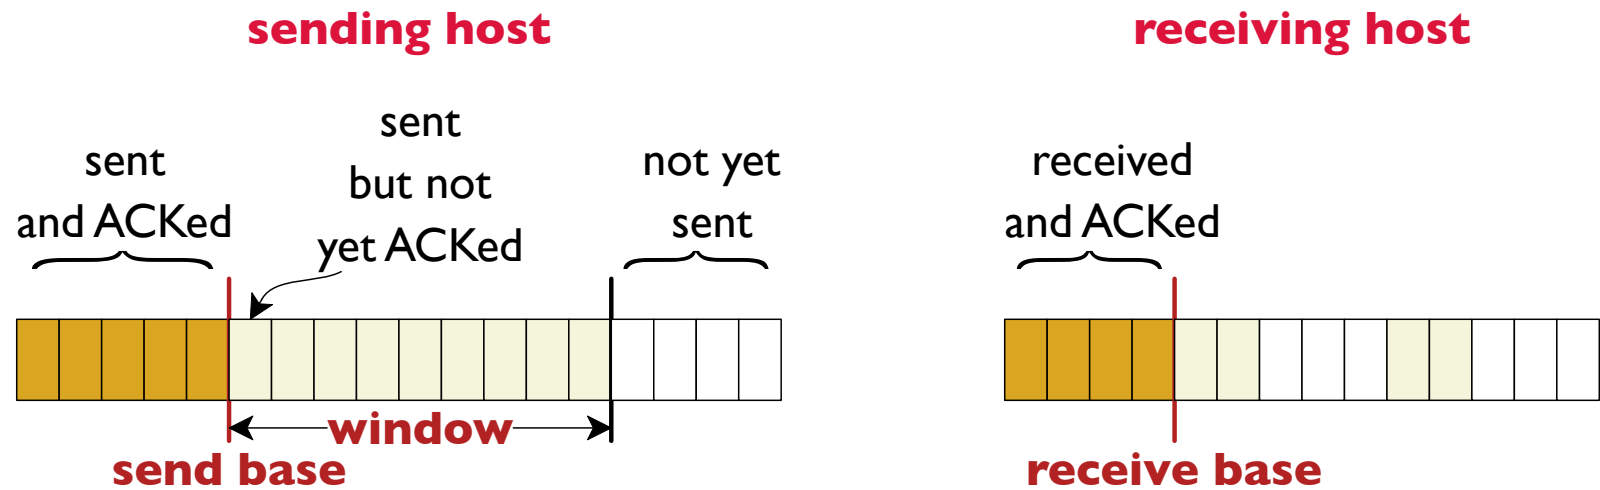
Go-Back-N : on timeout, re-send in window

Buffers



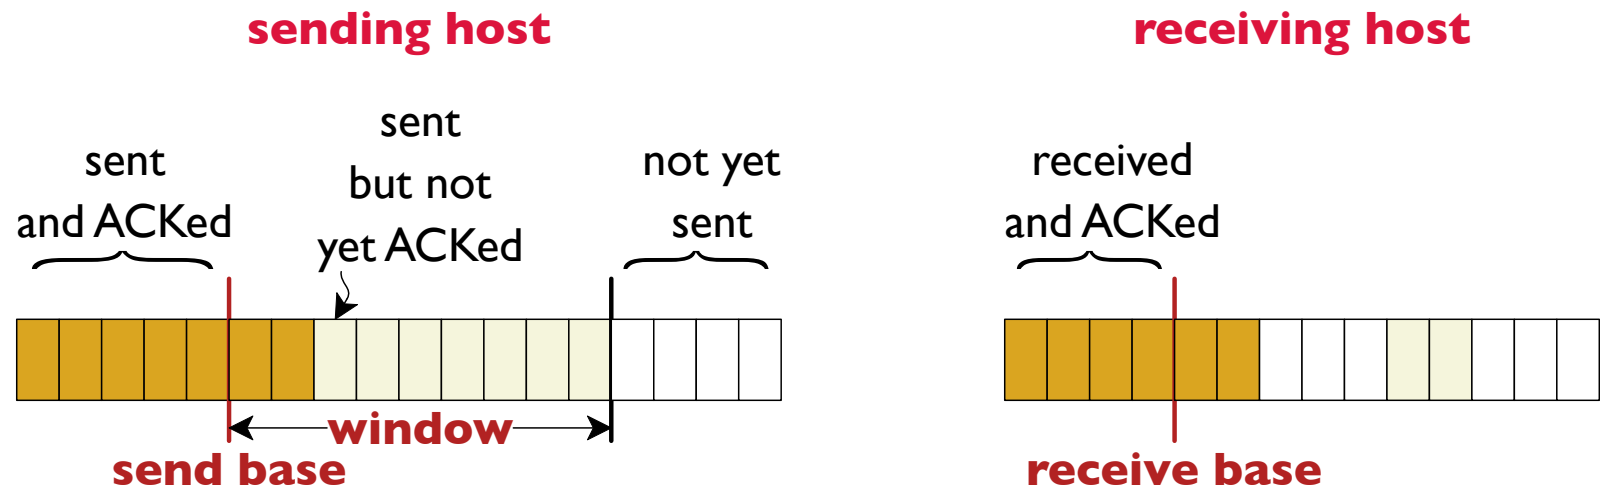
Go-Back-N : on timeout, re-send in window
Can use a **cumulative** ACK

Buffers



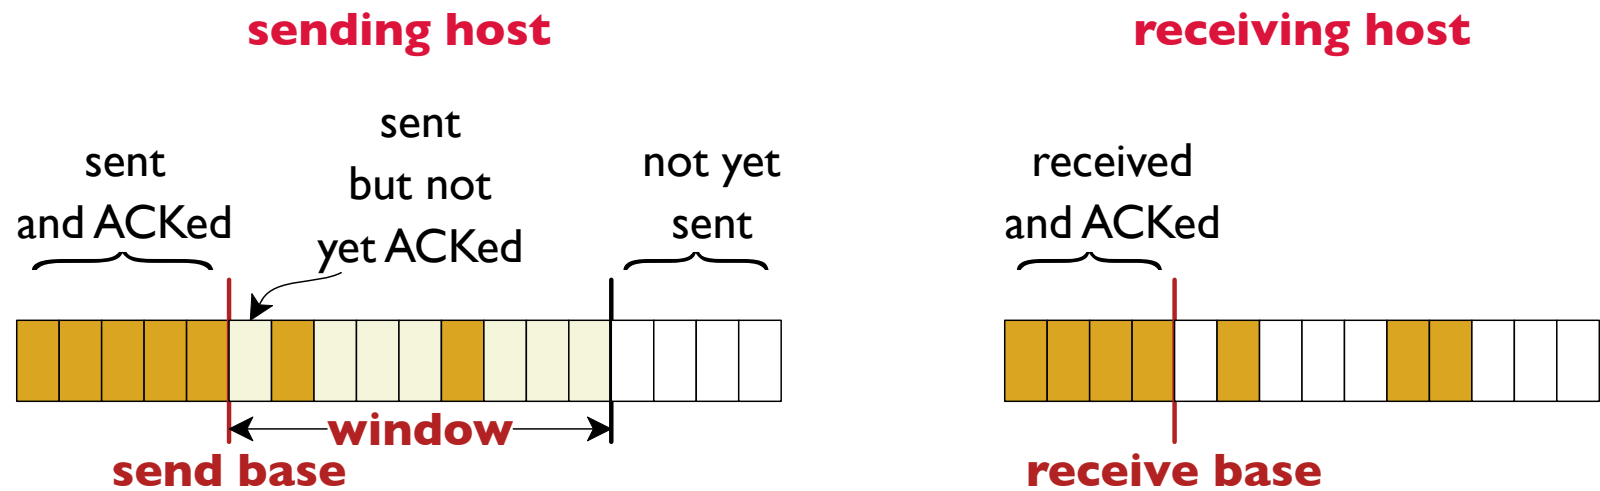
Go-Back-N : on timeout, re-send in window
Can use a **cumulative** ACK

Buffers



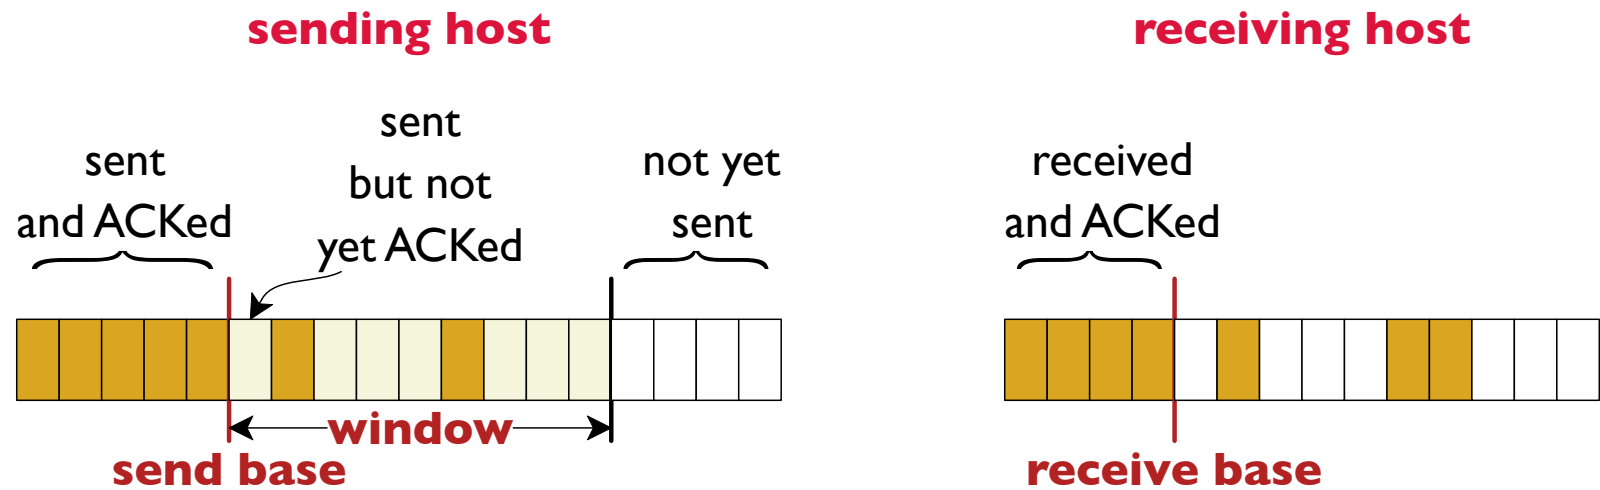
Go-Back-N : on timeout, re-send in window
Can use a **cumulative** ACK

Buffers



Selective repeat : on timeout, re-send unACKed

Buffers



Selective repeat : on timeout, re-send unACKed
Each packet must be specifically ACKed